

STAT 578 - Advanced Bayesian Modeling - Fall 2019

Assignment 4

Xiaoming Ji

Solution for Problem (a)

(i)


According to the formula for Moore's law and take log on T, we get:

$$\log(T) \approx \log(C * 2^{A/2}) \rightarrow \log(T) \approx \log(C) + 0.3465736 * A$$

This is the form of linear regression that make logT as response variable and A as explanatory variable. $\log(C)$ is the intercept and coefficient of A is 0.3465736. Thus we can say logT is roughly follow a simple linear regression on A.

(ii)

```
moores_df = read.csv("mooreslawdata.csv", header=TRUE)
with(moores_df, plot(log(Transistors) ~ Year))
```



Solution for Problem (b)

(i)

We first use classical linear model to estimate the parameters.

```
df_lm <- list(y = log(moores_df$Transistors),
             x_centered = moores_df$Year - mean(moores_df$Year))

model_lm = lm(y ~ x_centered, data=df_lm)
summary(model_lm)
```

```
##
## Call:
## lm(formula = y ~ x_centered, data = df_lm)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.0598 -0.4443  0.0891  0.6052  2.1820
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  17.726825   0.074044  239.41  <2e-16 ***
## x_centered    0.342122   0.005436   62.93  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.9395 on 159 degrees of freedom
## Multiple R-squared:  0.9614, Adjusted R-squared:  0.9612
## F-statistic: 3960 on 1 and 159 DF, p-value: < 2.2e-16
```

- β estimates are about 0.3 to 18, we choose to set initial β values at ± 200
- Regression error variance σ^2 estimate is about $(0.94^2) \approx 0.9$, we choose to set initial σ^2 values of 0.01 and 100

```
model {
  x_bar <- mean(x)
  for (i in 1:length(y)) {
    y[i] ~ dnorm(beta1 + beta2* (x[i] - x_bar), sigmasqinv)
  }

  beta1 ~ dnorm(0, 1e-06)
  beta2 ~ dnorm(0, 1e-06)
  sigmasqinv ~ dgamma(0.001, 0.001)
  sigmasq <- 1/sigmasqinv
}
```

```
library(rjags)
```

```
df_jags <- list(y = log(moores_df$Transistors),
               x = moores_df$Year)

initial_vals <- list(list(beta1 = 200, beta2 = 200, sigmasqinv = 100),
                    list(beta1 = 200, beta2 = -200, sigmasqinv = 100),
                    list(beta1 = -200, beta2 = 200, sigmasqinv = 0.01),
                    list(beta1 = -200, beta2 = -200, sigmasqinv = 0.01))
```

```

moores_model <- jags.model("moores.bug", df_jags, initial_vals, n.chains = 4)
update(moores_model, 1000)
coef_sample <- coda.samples(moores_model, c("beta1", "beta2", "sigmasq"), n.iter=2000)

gelman.diag(coef_sample, autoburnin=FALSE)

```

```

## Potential scale reduction factors:
##
##           Point est. Upper C.I.
## beta1           1           1
## beta2           1           1
## sigmasq         1           1
##
## Multivariate psrf
##
## 1

```

Gelman-Rubin statistic for β_1 , β_2 and σ^2 are all 1, thus we can declare convergence of them.

(ii)

```

summary(coef_sample)

##
## Iterations = 1001:3000
## Thinning interval = 1
## Number of chains = 4
## Sample size per chain = 2000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean          SD Naive SE Time-series SE
## beta1    17.7271 0.074781 0.0008361      8.362e-04
## beta2     0.3421 0.005456 0.0000610      6.014e-05
## sigmasq   0.8924 0.099704 0.0011147      1.111e-03
##
## 2. Quantiles for each variable:
##
##           2.5%      25%      50%      75%     97.5%
## beta1    17.5780 17.6780 17.7269 17.7779 17.8729
## beta2     0.3313  0.3385  0.3422  0.3457  0.3527
## sigmasq   0.7193  0.8231  0.8833  0.9550  1.1128

```

(iii)

```

beta2_sample = as.matrix(coef_sample)[,"beta2"]

```

Mean of slope,

```

mean(beta2_sample)

```

```

## [1] 0.3420985

```

95% posterior credible interval of slope,

```
quantile(beta2_sample, c(0.025, 0.975))
```

```
##      2.5%      97.5%  
## 0.331320 0.352727
```

The interval contains value (0.3465736) determined in part (a).

(iv)

```
beta1_sample = as.matrix(coef_sample)[,"beta1"]
```

Mean of intercept,

```
mean(beta1_sample)
```

```
## [1] 17.72707
```

95% posterior credible interval of intercept,

```
quantile(beta1_sample, c(0.025, 0.975))
```

```
##      2.5%      97.5%  
## 17.57802 17.87290
```

Solution for Problem (c)

(i)

```
model {  
  x_bar <- mean(x)  
  for (i in 1:length(y)) {  
    y[i] ~ dnorm(beta1 + beta2*(x[i] - x_bar), sigma_sqinv)  
  }  
  beta1 ~ dnorm(0, 1e-06)  
  beta2 ~ dnorm(0, 1e-06)  
  sigma_sqinv ~ dgamma(0.001, 0.001)  
  y_predict ~ dnorm(beta1 + beta2*(year_predict - x_bar), sigma_sqinv)  
  year_start <- x_bar - (beta1 / beta2)  
}
```

```
predict_model <- jags.model("moores_predict.bug",  
                           c(as.list(df_jags), year_predict = 2020),  
                           initial_vals, n.chains = 4)  
update(predict_model, 1000)  
predict_sample <- coda.samples(predict_model, c("y_predict", "year_start"), n.iter=2000)
```

```
gelman.diag(predict_sample, autoburnin=FALSE)
```

```
## Potential scale reduction factors:
```

```
##
```

```
##           Point est. Upper C.I.
```

```
## y_predict           1           1
```

```
## year_start          1           1
```

```
##
```

```
## Multivariate psrf
##
## 1
```

Gelman-Rubin statistic for the prediction is 1, thus we can declare convergence of them.

(ii)

```
summary(predict_sample)

##
## Iterations = 1001:3000
## Thinning interval = 1
## Number of chains = 4
## Sample size per chain = 2000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##              Mean      SD Naive SE Time-series SE
## y_predict      23.84 0.9490  0.01061      0.010571
## year_start 1950.25 0.8586  0.00960      0.009488
##
## 2. Quantiles for each variable:
##
##              2.5%    25%    50%    75%    97.5%
## y_predict      22    23.2  23.84  24.49  25.72
## year_start 1949 1949.7 1950.26 1950.85 1951.89
```

(iii)

95% posterior predictive interval for the transistor count, in billions,

```
exp(quantile(as.matrix(predict_sample)[,"y_predict"], c(0.025, 0.975))) / (10^9)

##          2.5%          97.5%
## 3.588029 147.620426
```

(iv)

To explain, we can assume the year when transistor is invented is the year for transistor count equal to 1. This means given our model,

$$\log(T) = \beta_1 + \beta_2 * (A_i - \bar{A})$$

We can set $T = 1$ to derive A_i as the year when transistor is invented. Therefore, we have,

$$A_i = \bar{A} - \beta_1 / \beta_2$$

Using the classical linear model, we can estimate this year as,

```
round(mean(moores_df$Year) - coef(model_lm)[[1]] / coef(model_lm)[[2]])

## [1] 1950
```

The 95% posterior interval for this quantity.

```
round(quantile(as.matrix(predict_sample)[, "year_start"], c(0.025, 0.975)))
```

```
## 2.5% 97.5%  
## 1949 1952
```

Note: According to Wikipedia, the first working device to be built was a point-contact transistor invented in 1947 by American physicists John Bardeen, Walter Brattain, and William Shockley at Bell Labs. Our estimates is very close.

Solution for Problem (d)

(i)

```
X <- model.matrix(model_lm)  
Nsim <- length(beta1_sample)  
error_sim <- matrix(NA, Nsim, nrow(X))  
post_beta <- as.matrix(coef_sample)[, c("beta1", "beta2")]  
post_simgasq <- as.matrix(coef_sample)[, c("sigmasq")]  
  
for(s in 1:Nsim) {  
  error_sim[s,] <- df_jags$y - X %*% cbind(post_beta[s,])  
}
```

(ii)


```
error_rep <- matrix(NA, Nsim, nrow(X))  
  
for(s in 1:Nsim) {  
  error_rep[s,] <- rnorm(nrow(X), 0, sqrt(post_simgasq[s]))  
}
```

(iii)

```
error_sim_std <- matrix(NA, Nsim, nrow(X))  
error_rep_std <- matrix(NA, Nsim, nrow(X))  
  
for(s in 1:Nsim) {  
  error_sim_std[s,] = error_sim[s,] / sqrt(post_simgasq[s])  
  error_rep_std[s,] = error_rep[s,] / sqrt(post_simgasq[s])  
}  
  
T <- apply(abs(error_sim_std), 1, max)  
T_Rep <- apply(abs(error_rep_std), 1, max)
```

(iv)

```
plot(T_Rep ~ T, pch=".", cex=2)  
abline(a=0, b=1)
```



(v)

```
mean(T_Rep >= T)
```

```
## [1] 0
```

p-value is 0 and this is the strong evidence for an outlier.

(vi)

The most extreme outlier is,

```
moores_df[unique(apply(abs(error_sim_std), 1, which.max)),]
```

```
## Processor Year Transistors
## 49 ARM 9TDMI 1999 111000
```

which is “ARM 9TDMI” produced in 1999 that only has 111,000 transistors. In this time period, other processors have tens of millions of transistors.