

ADVANCED BAYESIAN MODELING

EXAMPLE OF PRACTICAL MCMC:
RAT TUMOR DATA: MODEL 1

Rat Tumor Example

n_j = total number of rats in control group of experiment j

y_j = number in control group of experiment j that develop a tumor

θ_j = control group tumor probability in experiment j

j = $1, \dots, 71$

Data values for n and y (N and y) in file `rattumor.txt`

Consider model specified in JAGS code file `rattumor1.bug`:

```
model {  
  
  for (j in 1:length(y)) {  
    y[j] ~ dbin(theta[j], N[j])  
    theta[j] ~ dbeta(alpha, beta)  
  }  
  
  alpha ~ dexp(0.001)  
  beta ~ dexp(0.001)  
  
}
```

(Same as analyzed earlier)

This time, instead of using JAGS/rjags default options, we will

- ▶ use 4 chains
- ▶ initialize hyperparameters (α and β) to overdispersed combinations (different for each chain)
- ▶ run 1000 iterations for adaptation
- ▶ then run 1000 more iterations for burn-in
- ▶ finally run 2000 more iterations to check for convergence

Helpful: Determine crude preliminary values for parameters.

These will help determine what values could be considered “overdispersed.”

Since the θ s are probabilities, naive estimates are the empirical proportions

$$\hat{\theta}_j = \frac{y_j}{n_j}$$

The θ s have a $\text{Beta}(\alpha, \beta)$ distribution under the prior. In particular (BDA3, Table A.1):

$$\text{E}(\theta \mid \alpha, \beta) = \frac{\alpha}{\alpha + \beta} \qquad \text{var}(\theta \mid \alpha, \beta) = \frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)}$$

Solving these for α and β and then substituting the sample mean and variance of the $\hat{\theta}_j$ s for $\text{E}(\theta \mid \alpha, \beta)$ and $\text{var}(\theta \mid \alpha, \beta)$ gives ...

```
> d <- read.table("rattumor.txt", header=TRUE)

> thetahat <- d$y/d$N

> Etheta <- mean(thetahat)

> Vtheta <- var(thetahat)

> ( alphahat <- Etheta * (Etheta * (1 - Etheta) / Vtheta - 1) )
[1] 1.373948

> ( betahat <- alphahat * (1/Etheta - 1) )
[1] 8.573899
```

These values will help to initialize the chains.

Recall: Initial values should be overdispersed – extremely large or small.

Hyperparameters α and β must be positive, so let's choose initial values about 100 times larger and smaller than we expect.

For the 4 chains:

	Initial α	Initial β
Chain 1:	0.01	0.1
Chain 2:	100	0.1
Chain 3:	0.01	1000
Chain 4:	100	1000

Set up JAGS model in R with 4 chains initialized, and perform 1000 iterations of adaptation:

```
> library(rjags)  # automatically loads coda package
...
```

```
> initial.vals <- list(list(alpha=0.01, beta=0.1),
+                       list(alpha=100,  beta=0.1),
+                       list(alpha=0.01, beta=1000),
+                       list(alpha=100,  beta=1000))
```

```
> m1 <- jags.model("rattumor1.bug", d, initial.vals, n.chains=4, n.adapt=1000)
...
```


Perform 1000 iterations of burn-in, collect α and β for 2000 more iterations, and plot results:

```
> update(m1, 1000) # burn-in
|*****| 100%


> x1 <- coda.samples(m1, c("alpha","beta"), n.iter=2000)
|*****| 100%

> plot(x1, smooth=FALSE)
```


Trace of alpha




Density of alpha



Trace of beta



Density of beta




Chains are far from convergence, even after 4000 iterations.

Could run more iterations (try yourself), but don't expect much improvement.

For a Gibbs sampler, slow convergence might be due to high posterior correlations:

```
> plot(as.matrix(x1)[,c("alpha","beta")], pch=".", cex=2)
```



Results call into doubt earlier (default) analysis: There are regions of possibly high probability that were not sampled.

Fundamental problem: Priors were chosen close to flat for α and β , but using actual flat priors would give an improper posterior.

Better to try a more well-behaved prior ...