

# ADVANCED BAYESIAN MODELING

EXAMPLE OF PRACTICAL MCMC:  
**RAT TUMOR DATA: MODEL 2**

Consider alternative model specified in JAGS code file rattumor2.bug:

```
model {  
  
  for (j in 1:length(y)) {  
    y[j] ~ dbin(theta[j], N[j])  
    theta[j] ~ dbeta(alpha, beta)  
  }  
  
  alpha <- phi1 / phi2^2  
  beta <- (1-phi1) / phi2^2  
  
  phi1 ~ dunif(0,1)  
  phi2 ~ dunif(0,1000)  
  
}
```

(Also analyzed earlier)

Nodes `alpha` and `beta` are deterministic, so JAGS does not allow them to be initialized.

Instead initialize `phi1` and `phi2`, which are

$$\phi_1 = \frac{\alpha}{\alpha + \beta} \in (0, 1) \qquad \phi_2 = (\alpha + \beta)^{-1/2} \in (0, \infty)$$

Helpful preliminary values can be obtained using the preliminary  $\alpha$  and  $\beta$  values computed previously ...

```
> ( phi1hat <- alphahat / (alphahat + betahat) )
```

```
[1] 0.1381151
```

```
> ( phi2hat <- (alphahat + betahat)^(-1/2) )
```

```
[1] 0.3170556
```

Again, let's use 4 chains, initialized in a manner overdispersed relative to the preliminary values:

	Initial $\phi_1$	Initial $\phi_2$
Chain 1:	0.001	0.003
Chain 2:	0.9	0.003
Chain 3:	0.001	30
Chain 4:	0.9	30

Set up JAGS model in R with 4 chains initialized, and perform 1000 iterations of adaptation:

```
> library(rjags)
...

> initial.vals <- list(list(phi1=0.001, phi2=0.003),
+                       list(phi1=0.9,   phi2=0.003),
+                       list(phi1=0.001, phi2=30),
+                       list(phi1=0.9,   phi2=30))

> m2 <- jags.model("rattumor2.bug", d, initial.vals, n.chains=4, n.adapt=1000)
...
```


Perform 1000 iterations of burn-in, collect  $\alpha$  and  $\beta$  for 2000 more iterations, and plot results:

```
> update(m2, 1000) # burn-in
|*****| 100%


> x2 <- coda.samples(m2, c("alpha","beta"), n.iter=2000)
|*****| 100%

> plot(x2, smooth=FALSE)
```


Trace of alpha




Density of alpha



Trace of beta



Density of beta





Trace plots look good. Check a version of the Gelman-Rubin statistic, also called a **potential scale reduction factor**:

```
> gelman.diag(x2, autoburnin=FALSE)
```

Potential scale reduction factors:


	Point est.	Upper C.I.
alpha	1	1.01
beta	1	1.01

Multivariate psrf


1

```
> gelman.plot(x2, autoburnin=FALSE)
```

**alpha**



**beta**




Gelman-Rubin statistics appear legitimately near 1 (not just there by chance).


Check autocorrelation plot (pooled over all chains):

```
> autocorr.plot(x2)
```

**alpha**



**beta**



Some high autocorrelations, but essentially zero by lag 15.

Check effective sample sizes (summed over all chains):

```
> effectiveSize(x2)
      alpha      beta
708.8611 723.9213
```

Both exceed the suggested minimum value of 400.

These diagnostics all suggest adequate convergence and sampling, but could have run further diagnostics on other parameters ( $\theta$ s) had they been saved.