

ADVANCED BAYESIAN MODELING

Rat Growth Data Example: JAGS Analysis

y_{ij} = weight of rat j at age x_i (days) $i = 1, \dots, 5$ $j = 1, \dots, 30$

$$x_1 = 8 \quad x_2 = 15 \quad x_3 = 22 \quad x_4 = 29 \quad x_5 = 36$$

To aid computation, use centered ages:

$$x_i^{\text{cent}} = x_i - \bar{x}$$

```
> ratgrowth <- read.csv("ratgrowth.csv", header=TRUE)
```

```
> head(ratgrowth)
```

	day8	day15	day22	day29	day36
1	151	199	246	283	320
2	145	199	249	293	354
3	147	214	263	312	328
4	155	200	237	272	297
5	135	188	230	280	323
6	159	210	252	298	331

Note: Index j corresponds to rows, index i to columns.

We fit normal-theory model

$$y_{ij} \mid \beta^{(j)}, \sigma_y^2, x_i \sim \text{indep. N}(\beta_1^{(j)} + \beta_2^{(j)} x_i^{\text{cent}}, \sigma_y^2)$$

$$\beta^{(j)} \mid \mu_\beta, \Sigma_\beta \sim \text{iid N}(\mu_\beta, \Sigma_\beta)$$

$$\mu_\beta = \begin{pmatrix} \mu_{\beta_1} \\ \mu_{\beta_2} \end{pmatrix} \quad \Sigma_\beta = \begin{pmatrix} \sigma_{\beta_1}^2 & \rho \sigma_{\beta_1} \sigma_{\beta_2} \\ \rho \sigma_{\beta_1} \sigma_{\beta_2} & \sigma_{\beta_2}^2 \end{pmatrix}$$

Parameters of interest include σ_y^2 , μ_{β_1} , μ_{β_2} , $\sigma_{\beta_1}^2$, $\sigma_{\beta_2}^2$, and ρ .

Using centered predictors implies that the intercept $\beta_1^{(j)}$ represents predicted weight at the average age ($\bar{x} = 22$).

Simple least squares analyses can guide choice of priors and initialization:

```
> betahat <- matrix(NA, nrow(ratgrowth), 2)
> for(j in 1:nrow(ratgrowth))
+   betahat[j,] <- lsfit(rbind(8, 15, 22, 29, 36) - 22, t(ratgrowth[j,]))$coef

> apply(betahat, 2, mean)
[1] 242.653333    6.185714

> var(betahat)
      [,1]      [,2]
[1,] 210.399816  4.8615764
[2,]   4.861576  0.3344968
```

We need to choose proper priors:

$$\sigma_y^2 \sim \text{Inv-gamma}(0.0001, 0.0001)$$

$$\mu_\beta \sim \text{N}(0, 1000^2 I)$$

$$\Sigma_\beta^{-1} \sim \text{Wishart}_2(\Sigma_0^{-1}/2)$$

To ensure scaling is not grossly incorrect, let

$$\Sigma_0 = \begin{pmatrix} 100 & 0 \\ 0 & 0.1 \end{pmatrix}$$

based on preliminary least squares results.

```

data {
  dimY <- dim(weight)
  agecent <- age - mean(age)
}
model {
  for (j in 1:dimY[1]) {
    for (i in 1:dimY[2]) {
      weight[j,i] ~ dnorm(beta[1,j] + beta[2,j]*agecent[i], sigmasqyinv)
    }
    beta[1:2,j] ~ dmnorm(mubeta, Sigmabetainv)
  }

  mubeta ~ dmnorm(mubeta0, Sigmamubetainv)
  Sigmabetainv ~ dwish(2*Sigma0, 2)
  sigmasqyinv ~ dgamma(0.0001, 0.0001)

  Sigmabeta <- inverse(Sigmabetainv)
  rho <- Sigmabeta[1,2] / sqrt(Sigmabeta[1,1] * Sigmabeta[2,2])
  sigmasqy <- 1/sigmasqyinv
}

```


Note:

- ▶ Age centering performed in data block
- ▶ Explicit row range in `beta[1:2,j]` needed for JAGS to determine dimension
- ▶ `inverse` computes matrix inverse

Set up data and constants:

```
d1 <- list(weight = ratgrowth,  
           age = c(8, 15, 22, 29, 36),  
           mubeta0 = c(0, 0),  
           Sigmamubetainv = rbind(c(0.000001, 0),  
                                   c(0, 0.000001)),  
           Sigma0 = rbind(c(100, 0),  
                           c(0, 0.1)))
```

Set up initializations (extreme relative to data) for 4 chains:

```
inits1 <- list(list(sigmatqyinv = 10, mubeta = c(1000, 1000),  
                  Sigmabetainv = rbind(c(100, 0),  
                                       c(0, 100))),  
              list(sigmatqyinv = 0.001, mubeta = c(-1000, 1000),  
                  Sigmabetainv = rbind(c(100, 0),  
                                       c(0, 100))),  
              list(sigmatqyinv = 10, mubeta = c(1000, -1000),  
                  Sigmabetainv = rbind(c(0.001, 0),  
                                       c(0, 0.001))),  
              list(sigmatqyinv = 0.001, mubeta = c(-1000, -1000),  
                  Sigmabetainv = rbind(c(0.001, 0),  
                                       c(0, 0.001))))
```

Note: Only top-level nodes are initialized.

```

> library(rjags)

> m1 <- jags.model("ratgrowth1.bug", d1, inits1, n.chains=4, n.adapt=1000)
...

> update(m1, 1000)  # burn-in
|*****| 100%

> x1 <- coda.samples(m1, c("mubeta","Sigmabeta","sigmasqy"), n.iter=2000)
|*****| 100%

```

Note: Monitoring the full matrix Sigmabeta (including both off-diagonal elements)

```
> gelman.diag(x1, autoburnin=FALSE, multivariate=FALSE)
```

Potential scale reduction factors:

	Point est.	Upper C.I.
Sigmabeta[1,1]	1.31	2.03
Sigmabeta[2,1]	1.30	1.67
Sigmabeta[1,2]	1.30	1.67
Sigmabeta[2,2]	1.31	2.09
mubeta[1]	2.39	12.94
mubeta[2]	31.95	199.76
sigmasqy	7.92	48.96

Clearly requires more burn-in.

Remark: multivariate=FALSE because of a singularity issue

Chains eventually appear to converge by 31000 iterations (not shown).

Now compute the iterations to be used for inference:

```
> x1 <- coda.samples(m1, c("beta","mubeta","Sigmabeta","sigmasqy","rho"),
+                     n.iter=2000)
|*****| 100%

> effectiveSize(x1[,c("mubeta[1]","mubeta[2]","Sigmabeta[1,1]","Sigmabeta[1,2]","
+                     "Sigmabeta[2,2]","sigmasqy","rho")])
```

mubeta[1]	mubeta[2]	Sigmabeta[1,1]	Sigmabeta[1,2]	Sigmabeta[2,2]
7528.876	5263.405	6756.783	4983.577	3174.161
sigmasqy	rho			
2809.634	3059.025			

```
> summary(x1[,c("mubeta[1]","mubeta[2]","Sigmabeta[1,1]","Sigmabeta[1,2]",
+               "Sigmabeta[2,2]","sigmasqy","rho")])
```

Iterations = 31001:33000

Thinning interval = 1

Number of chains = 4

Sample size per chain = 2000

1. Empirical mean and standard deviation for each variable,
plus standard error of the mean:

	Mean	SD	Naive SE	Time-series SE
mubeta[1]	242.6663	2.74148	0.030651	0.031593
mubeta[2]	6.1867	0.10715	0.001198	0.001478
Sigmabeta[1,1]	217.3919	61.70650	0.689900	0.751189
Sigmabeta[1,2]	4.9700	1.96815	0.022005	0.027941
Sigmabeta[2,2]	0.2694	0.09612	0.001075	0.001707
sigmasqy	37.6433	5.87822	0.065720	0.111669
rho	0.6478	0.13426	0.001501	0.002431





2. Quantiles for each variable:

	2.5%	25%	50%	75%	97.5%
mubeta[1]	237.2535	240.8789	242.6421	244.4544	248.0520
mubeta[2]	5.9758	6.1160	6.1865	6.2576	6.3993
Sigmabeta[1,1]	127.9961	173.6677	207.4945	250.1384	365.0139
Sigmabeta[1,2]	1.9920	3.6102	4.6657	6.0064	9.6293
Sigmabeta[2,2]	0.1276	0.2025	0.2550	0.3186	0.5031
sigmasqy	28.0428	33.4059	37.0914	41.1130	50.5589
rho	0.3369	0.5699	0.6658	0.7464	0.8567


Note: ρ is almost certainly positive.

Plot estimates of marginal posterior densities:

```
> par(mfrow=c(2,2))  
> densplot(x1[,c("mubeta[1]", "mubeta[2]", "Sigmabeta[1,1]", "Sigmabeta[1,2]",  
+               "Sigmabeta[2,2]", "sigmasqy", "rho")])
```




Density of Sigmabeta[2,2]




N = 2000 Bandwidth = 0.01521

Density of sigmasqy



N = 2000 Bandwidth = 1.01

Density of rho



N = 2000 Bandwidth = 0.02315

Compute JAGS version of DIC:

```
> dic.samples(m1,2000)
|*****| 100%
Mean deviance: 967.7
penalty 51.68
Penalized deviance: 1019
```

Effective number of parameters: about 52

Actual number of parameters: $60 (\beta^{(j)}\text{s}) + 1 (\sigma_y^2) + 2 (\mu_\beta) + 3 (\Sigma_\beta)$

Alternative JAGS model using a scaled inverse Wishart:

```
data {  
  dimY <- dim(weight)  
  agecent <- age - mean(age)  
}  
model {  
  for (j in 1:dimY[1]) {  
    for (i in 1:dimY[2]) {  
      weight[j,i] ~ dnorm(beta[1,j] + beta[2,j]*agecent[i], sigmasqyinv)  
    }  
    beta[1:2,j] <- mubeta + xi * eta[,j]  
    eta[1:2,j] ~ dmnorm(zeros2, Sigmaetainv)  
  }  
  xi[1] ~ dunif(0,1000)  
  xi[2] ~ dunif(0,1000)
```

(continued...)

```

mubeta ~ dmnorm(mubeta0, Sigmamubetainv)
Sigmaetainv ~ dwish(identity2x2, 3)
sigmasqyinv ~ dgamma(0.0001, 0.0001)

Sigmaeta <- inverse(Sigmaetainv)
sigmasqbeta[1] <- xi[1]^2 * Sigmaeta[1,1]
sigmasqbeta[2] <- xi[2]^2 * Sigmaeta[2,2]
rho <- Sigmaeta[1,2] / sqrt(Sigmaeta[1,1] * Sigmaeta[2,2])
sigmasqy <- 1/sigmasqyinv
}

```

Constants zeros2 and identity2x2 are specified with the data.

Remark: Convergence requires more iterations.